

# Introduction to **Information Retrieval**


Lecture 18: Link analysis

# Today's lecture

---

- Anchor text
- Link analysis for ranking
  - Pagerank and variants

# The Web as a Directed Graph



**Assumption 1:** A hyperlink between pages denotes author perceived relevance (quality signal)


**Assumption 2:** The text in the anchor of the hyperlink describes the target page (textual context)

# Anchor Text

## *WWW Worm* - McBryan [Mcbr94]

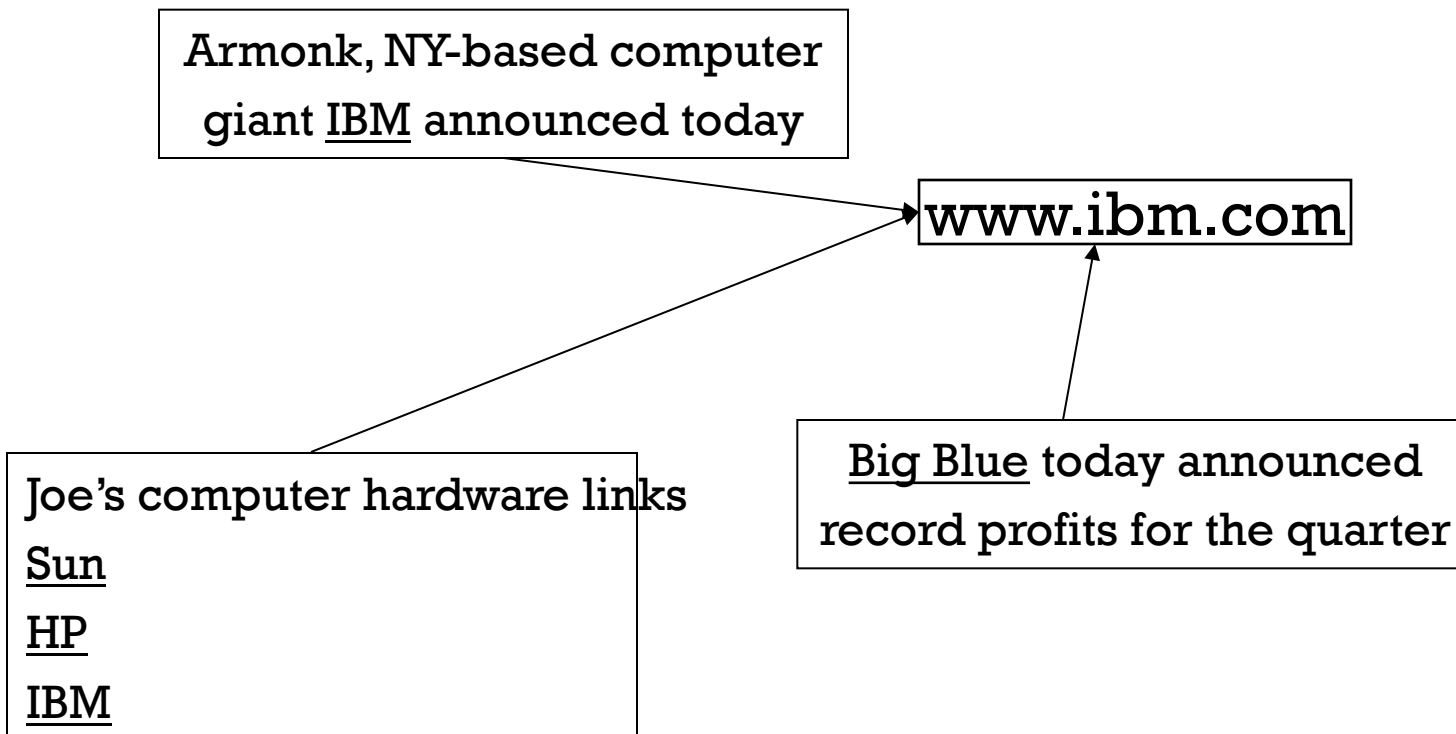
---

- For **ibm** how to distinguish between:
  - IBM's home page (mostly graphical)
  - IBM's copyright page (high term freq. for 'ibm')
  - Rival's spam page (arbitrarily high term freq.)



# Indexing anchor text

- When indexing a document  $D$ , include anchor text from links pointing to  $D$ .



# Indexing anchor text

---

- Can sometimes have unexpected side effects - *e.g., evil empire.*
- Can score anchor text with weight depending on the authority of the anchor page's website
  - E.g., if we were to assume that content from cnn.com or yahoo.com is authoritative, then trust the anchor text from them

# Anchor Text

---

- Other applications
  - Weighting/filtering links in the graph
  - Generating page descriptions from anchor text

# Citation Analysis


---

- Citation frequency
- Co-citation coupling frequency
  - Cocitations with a given author measures “impact”
  - Cocitation analysis
- Bibliographic coupling frequency
  - Articles that co-cite the same articles are related
- Citation indexing
  - Who is this author cited by? (Garfield 1972)
- Pagerank preview: Pinski and Narin '60s



# Query-independent ordering

- First generation: using link counts as simple measures of popularity.
- Two basic suggestions:
  - Undirected popularity:
    - Each page gets a score = the number of in-links plus the number of out-links ( $3+2=5$ ).
  - Directed popularity:
    - Score of a page = number of its in-links (3).



# Query processing

---

- First retrieve all pages meeting the text query (say ***venture capital***).
- Order these by their link popularity (either variant on the previous slide).
- More nuanced – use link counts as a measure of static goodness (Lecture 7), combined with text match score

# Spamming simple popularity

---

- *Exercise:* How do you spam each of the following heuristics so your page gets a high score?
  1. Each page gets a static score = the number of in-links plus the number of out-links.
  2. Static score of a page = number of its in-links.

# Ideas of Pagerank


---

- Inlinks as votes
  - [www.stanford.edu](http://www.stanford.edu) has 23,400 inlinks
  - [www.joe-schmoe.com](http://www.joe-schmoe.com) has 1 inlink
- Web pages are not equally “important”
  - [www.joe-schmoe.com](http://www.joe-schmoe.com) → p1
  - vs. [www.stanford.edu](http://www.stanford.edu) → p2
- Are all inlinks equal?
  - Recursive question!

# Pagerank scoring

- Imagine a browser doing a *random walk* on web pages:

- Start at a random page




- At each step, go out of the current page along one of the links on that page, equiprobably

- “In the steady state” each page has a long-term visit rate - use this as the page’s score.

# Example – the Simple “Flow” Model

The web in 1839



$$y = y/2 + a/2$$

$$a = y/2 + m$$

$$m = a/2$$


# Solving the flow equations

---

- 3 equations, 3 unknowns, no constants
  - No unique solution
  - All solutions equivalent modulo scale factor
- Additional constraint forces uniqueness
  - $y+a+m = 1$
  - $y = 2/5, a = 2/5, m = 1/5$
- Gaussian elimination method works for small examples, but we need a better method for large graphs

# Example – the Simple “Flow” Model

The web in 1839



$$y_{new} = y_{old} / 2 + a_{old} / 2$$

$$a_{new} = y_{old} / 2 + m_{old}$$

$$m_{new} = a_{old} / 2$$

$$y = 1/3 \quad 1/3 \quad 5/12 \quad \dots \quad 2/5$$

$$a = 1/3 \quad 1/2 \quad 1/3 \quad \dots \quad 2/5$$



$$m = 1/3 \quad 1/6 \quad 1/4 \quad \dots \quad 1/5$$

Matrix-based characterization of the computation is simpler and more useful for the general case.




# Markov chains

- A Markov chain consists of  $n$  states, plus an  $n \times n$  transition probability matrix  $\mathbf{P}$ .
- **At each step, we are in exactly one of the states.**
- For  $1 \leq i, j \leq n$ , the matrix entry  $P_{ij}$  tells us the probability of  $j$  being the next state, given we are currently in state  $i$ .



# Markov chains


- Clearly, for all  $i$ ,  $\sum_{j=1}^n P_{ij} = 1$ .
- Markov chains are abstractions of random walks.



	y	a	m
y	1/2	1/2	0
a	1/2	0	1/2
m	0	1	0

# Ergodic Markov chains

- A Markov chain is ergodic if
  - you have a path from any state to any other
  - For any start state, after a finite transient time  $T_0$ , the probability of being in any state at a fixed time  $T > T_0$  is nonzero.



# Ergodic Markov chains

---

- For any ergodic Markov chain, there is a unique long-term visit rate for each state.
  - *Steady-state probability distribution.*
- Over a long time-period, we visit each state in proportion to this rate.
- It doesn't matter where we start.

# Probability vectors

---

- A probability (row) vector  $\mathbf{x} = (x_1, \dots, x_n)$  tells us where the walk is at any point.
- E.g.,  $(\underset{1}{000}\dots\underset{i}{1}\dots\underset{n}{000})$  means we're in state  $i$ .

More generally, the vector  $\mathbf{x} = (x_1, \dots, x_n)$  means the walk is in state  $i$  with probability  $x_i$ .

$$\sum_{i=1}^n x_i = 1.$$

# Change in probability vector

---

- If the probability vector is  $\mathbf{x} = (x_1, \dots, x_n)$  at this step, what is it at the next step?
- Recall that row  $i$  of the transition prob. Matrix  $\mathbf{P}$  tells us where we go next from state  $i$ .
- So from  $\mathbf{x}$ , our next state is distributed as  $\mathbf{xP}$ .

# How do we compute this vector?

---

- Let  $\mathbf{a} = (a_1, \dots, a_n)$  denote the row vector of steady-state probabilities.
- If our current position is described by  $\mathbf{a}$ , then the next step is distributed as  $\mathbf{aP}$ .
- But  $\mathbf{a}$  is the steady state, so  $\mathbf{a} = \mathbf{aP}$ .
- Solving this matrix equation gives us  $\mathbf{a}$ .
  - So  $\mathbf{a}$  is the (left) eigenvector for  $\mathbf{P}$ .
  - (Corresponds to the “principal” eigenvector of  $\mathbf{P}$  with the largest eigenvalue.)
  - Transition probability matrices always have largest eigenvalue 1.


# One way of computing $\mathbf{a}$

---

- Recall, regardless of where we start, we eventually reach the steady state  $\mathbf{a}$ .
- Start with any distribution (say  $\mathbf{x}=(1/n, 1/n, \dots, 1/n)$ ).
- After one step, we're at  $\mathbf{xP}$ ;
- after two steps at  $\mathbf{xP}^2$ , then  $\mathbf{xP}^3$  and so on.
- “Eventually” means for “large”  $k$ ,  $\mathbf{xP}^k = \mathbf{a}$ .
- Algorithm: multiply  $\mathbf{x}$  by increasing powers of  $\mathbf{P}$  until the product looks stable.



# Power Iteration Example



$\mathbf{x}_0 = (1/3, 1/3, 1/3)$

	<b>P</b>		
	y	a	m
y	1/2	1/2	0
a	1/2	0	1/2
m	0	1	0

$$y_{new} = y_{old} / 2 + a_{old} / 2$$
$$a_{new} = y_{old} / 2 + m_{old}$$
$$m_{new} = a_{old} / 2$$


		$\mathbf{x}_0$	$\mathbf{x}_1$	$\mathbf{x}_2$	$\mathbf{x}_3$		$\mathbf{x}_t$
y	=	1/3	1/3	5/12	3/8		2/5
a		1/3	1/2	1/3	11/24	...	2/5
m		1/3	1/6	1/4	1/6		1/5

# Spider traps

---

- A group of pages is a **spider trap** if there are no links from within the group to outside the group
  - Random surfer gets trapped
- Spider traps violate the conditions needed for the random walk theorem

# Microsoft becomes a spider trap



	y	a	m
y	1/2	1/2	0
a	1/2	0	1/2
m	0	0	1


y		1/3	1/3	1/4	5/24		0
a	=	1/3	1/6	1/6	1/8	...	0
m		1/3	1/2	7/12	2/3		1

# Random teleports

---

- The Google solution for spider traps
- At each time step, the random surfer has two options:
  - With probability  $\beta$ , follow a link at random
  - With probability  $1-\beta$ , jump to some page uniformly at random
  - Common values for  $\beta$  are in the range 0.8 to 0.9
- Surfer will teleport out of spider trap within a few time steps


# Random teleports ( $\beta = 0.8$ )



$$0.8 \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1/2 \\ 0 & 0 & 1 \end{bmatrix} + 0.2 \begin{bmatrix} 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \end{bmatrix}$$

$$\begin{matrix} y \\ a \\ m \end{matrix} \begin{bmatrix} 7/15 & 7/15 & 1/15 \\ 7/15 & 1/15 & 7/15 \\ 1/15 & 1/15 & 13/15 \end{bmatrix}$$

# Random teleports ( $\beta = 0.8$ )



$$0.8 \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 0 \\ 0 & 1/2 & 1 \end{bmatrix} + 0.2 \begin{bmatrix} 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \end{bmatrix}$$

y	7/15	7/15	1/15
a	7/15	1/15	7/15
m	1/15	1/15	13/15


y		0.33	0.33	0.28	0.26		7/33
a	=	0.33	0.20	0.20	0.18	...	5/33
m		0.33	0.47	0.52	0.56		7/11

# Dead ends

---

- Pages with no outlinks are “dead ends” for the random surfer
  - Nowhere to go on next step
- Especially common for Web Search Engines
  - URLs that have not yet been crawled

# Microsoft becomes a dead end



$$0.8 \begin{bmatrix} 1/2 & 1/2 & 0 \\ 1/2 & 0 & 1/2 \\ 0 & 0 & 0 \end{bmatrix}$$

$$+ 0.2 \begin{bmatrix} 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \\ 1/3 & 1/3 & 1/3 \end{bmatrix}$$

$$\begin{matrix} y \\ a \\ m \end{matrix} \begin{bmatrix} 7/15 & 7/15 & 1/15 \\ 7/15 & 1/15 & 7/15 \\ 1/15 & 1/15 & 1/15 \end{bmatrix}$$

Non-stochastic!

$$\begin{matrix} y \\ a \\ m \end{matrix} = \begin{bmatrix} 0.33 & 0.33 & 0.262 & 0.216 & & 0 \\ 0.33 & 0.20 & 0.182 & 0.143 & \dots & 0 \\ 0.33 & 0.20 & 0.129 & 0.111 & & 0 \end{bmatrix}$$



# Dealing with dead-ends

---

- Teleport
  - Follow random teleport links **with probability 1.0** from dead-ends
  - Adjust matrix **accordingly**
    - **How?**
- (Suggested by Google) prune and propagate
  - Preprocess the graph to eliminate dead-ends
  - Might require multiple passes
  - Compute page rank on reduced graph
  - **Approximate** values for deadends by propagating values from reduced graph

Q: Why approximate values and why errors are insignificant?

# Pagerank summary

---

- Preprocessing:

- Given graph of links, build matrix  $\mathbf{P}$ .
- From it compute  $\mathbf{a}$ .
  - $\mathbf{a}$  is the principle eigen vector of a matrix  $\tilde{\mathbf{P}}$

$$\tilde{\mathbf{P}} = (1 - \beta)\mathbf{P} + \beta\mathbf{T}, \quad \mathbf{T}_{i,j} = \frac{1}{n}$$

- The entry  $a_i$  is a number between 0 and 1: the pagerank of page  $i$ .

- Query processing:

- Retrieve pages meeting query.
- Rank them by their pagerank.
- Order is *query-independent*.

# The reality

---

- Pagerank is used in google, but is hardly the full story of ranking
  - Many sophisticated features are used
  - Some address specific query classes
  - Machine learned ranking (Lecture 15) heavily used
- Pagerank still very useful for things like crawl policy

# Pagerank: Issues and Variants

---

- How realistic is the random surfer model?
  - (Does it matter?)
  - What if we modeled the back button?
  - Surfer behavior sharply skewed towards short paths
  - Search engines, bookmarks & directories make jumps non-random.
- Biased Surfer Models
  - Weight edge traversal probabilities based on match with topic/query (non-uniform edge selection)
  - Bias jumps to pages on topic (e.g., based on personal bookmarks & categories of interest)

# Topic Specific Pagerank

---

- Goal – pagerank values that depend on query *topic*
- Conceptually, we use a random surfer who teleports, with say 10% probability, using the following rule:
  - Selects a topic (say, one of the 16 top level ODP categories) based on a query & user -specific distribution over the categories
  - Teleport to a page uniformly at random within the chosen topic
- Sounds hard to implement: can't compute PageRank at query time!

# Topic Specific Pagerank

---

- **Offline:** Compute pagerank for *individual* topics
  - Query independent as before
  - Each page has multiple pagerank scores – one for each ODP category, with teleportation only to that category
- **Online:** Query context classified into (distribution of weights over) topics
  - Generate a dynamic pagerank score for each page – weighted sum of topic-specific pageranks

# Influencing PageRank (“Personalization”)

- Input:

- Web graph  $W$
- Influence vector  $\mathbf{v}$  over topics  
 $\mathbf{v} : (\text{page} \rightarrow \text{degree of influence})$


Vector has one component for each topic

- Output:

- Rank vector  $\mathbf{r}$ : (page  $\rightarrow$  page importance wrt  $\mathbf{v}$ )

- $\mathbf{r} = \text{PR}(W, \mathbf{v})$

# Non-uniform Teleportation



Teleport with 10% probability to a Sports page



# Interpretation of Composite Score


---

- Given a set of personalization vectors  $\{\mathbf{v}_j\}$

$$\sum_j [w_j \cdot \text{PR}(W, \mathbf{v}_j)] = \text{PR}(W, \sum_j [w_j \cdot \mathbf{v}_j])$$


Given a user's preferences over topics, express as a combination of the “basis” vectors  $\mathbf{v}_j$

# Interpretation




10% Sports teleportation

# Interpretation



10% Health teleportation

# Interpretation



$pr = (0.9 PR_{\text{sports}} + 0.1 PR_{\text{health}})$  gives you:  
9% sports teleportation, 1% health teleportation

# Resources

---

- IIR Chap 21
- <http://www2004.org/proceedings/docs/1p309.pdf>
- <http://www2004.org/proceedings/docs/1p595.pdf>
- <http://www2003.org/cdrom/papers/refereed/p270/kamvar-270-xhtml/index.html>
- <http://www2003.org/cdrom/papers/refereed/p641/xhtml/p641-mccurley.html>