

# Neural Language Models

Wei Wang @ CSE, UNSW

November 13, 2019

- NLM by Bengio *et al.* (2003).
- NLM based on RNN.

# Language Model (LM)

- LM by definition:  $P(w_1, w_2, \dots, w_t)$ , where  $w_i \in V$ .
- Key to the LM:  $P(w_t | w_1, w_2, \dots, w_{t-1})$ ,  $\forall t$

$n$ -gram LM:

- $(n - 1)$ -th order Markov assumption
- memorize the MLE of the distribution
- smoothing

# Generalization for LM

Different (non-OOV) cases for generalization: (let  $s$  be a sequence of words)

- $s$  has been seen in the training data
- $s$  has never been seen in the training data, but part of it has
- $s$  has never been seen in the training data, and neither does any part of it

Examples:


- I saw Mike yesterday
- I saw Sabrunil yesterday
- Sid Meier's maagnum opus

**Question:** Why not just use  $P_{MLE}(s)$ ?

- Using  $k$ -th order Markov assumption; therefore, only need to model  $P(w_{+1} \mid w_{-k+1} \dots w_0)$ , or  $P(w_{+1} \mid \text{context})$ .
- Prototype design:  
 $\mathbf{y} = f(\mathbf{c}) = f([\mathbf{w}_{t-k+1} ; \mathbf{w}_{t-k+2} ; \dots ; \mathbf{w}_{t-1}])$ . Use shapes as cues:
  - $\mathbf{y}$ :
  - $\mathbf{c}$ :
  - Many reasons to employ a hidden layer  $\mathbf{h}$  to learned a better representation.
- Initial design:
  - $\mathbf{h} = \sigma(\mathbf{W}_{ch}\mathbf{c} + \mathbf{b}_h)$
  - $\mathbf{y} = \text{Softmax}(\mathbf{W}_{hy}\mathbf{h})$  (Why no bias term?)
- # of Parameters:
  - Note that  $\dim(w_i) = V$ , and let  $d = \dim(\mathbf{h})$
  - $\underbrace{(kV)d}_{\mathbf{W}_{ch}} + \underbrace{d}_{\mathbf{b}_h} + \underbrace{dV}_{\mathbf{W}_{hy}}$

- $\mathbf{W}_{ch}$  gives different embedding to the same word at different location in the context.
- 
- Share the embedding matrix (i.e., word embedding is independent of the location)
- Final design:
  - $\mathbf{c} = [\mathbf{E}\mathbf{w}_{-k+1} ; \mathbf{E}\mathbf{w}_{-k+2} ; \dots ; \mathbf{E}\mathbf{w}_0]$ , and  $\mathbf{w}_i$  is its one-hot encoding.
  - $\mathbf{h} = \sigma(\mathbf{W}_{ch}\mathbf{c} + \mathbf{b}_h)$
  - $\mathbf{y} = \text{Softmax}(\mathbf{W}_{hy}\mathbf{h})$
- # of Parameters:
  - Let  $m = \dim(\mathbf{E}\mathbf{w}_i)$ ,  $d = \dim(\mathbf{h})$
  - $\underbrace{Vm}_{\mathbf{E}} + \underbrace{kmd}_{\mathbf{W}_{ch}} + \underbrace{d}_{\mathbf{b}_h} + \underbrace{dV}_{\mathbf{W}_{hy}}$
- Can we get rid of the fixed order Markov assumption?

# Illustration



**Figure 7.13** learning all the way back to embeddings. notice that the embedding matrix  $E$  is shared among the 3 context words.

- Design:
  - We already have  $\mathbf{h}$  in RNN (at every timestamp)
    - $[\mathbf{h}_1, \mathbf{h}_2, \dots, \mathbf{h}_t] = \text{RNN}([\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_t])$
    - $\mathbf{y}_i = \text{Softmax}(\mathbf{W}_{hy} \mathbf{h}_i)$
- Designing the Loss function:
  - $\mathbf{t}_i = \text{onehot}(w_{+1})$
  - Applying the cross-entropy loss:  $\ell_i = -\log(\mathbf{y}_{w_{+1}})$
  - Average over all timestamps:  $L = -\frac{1}{T} \sum_{i=1}^T \ell_i$
  - Comment:
    - this is a common paradigm (applies to many models, such as Bengio's NLM, word2vec, etc.)
    - note that  $L$  is determined by the model parameters (not shown explicitly above)



## A RNN Language Model

output distribution

$$\hat{y}^{(t)} = \text{softmax}(Uh^{(t)} + b_2) \in \mathbb{R}^{|V|}$$

hidden states

$$h^{(t)} = \sigma(W_h h^{(t-1)} + W_e e^{(t)} + b_1)$$


$h^{(0)}$  is the initial hidden state

word embeddings

$$e^{(t)} = Ex^{(t)}$$

words / one-hot vectors

$$x^{(t)} \in \mathbb{R}^{|V|}$$



# Neural LM with a Cache

- Problem: temporal locality

Kookaburras are terrestrial tree kingfishers of the genus *Dacelo* native to Australia and New Guinea, ...  
The genus \_\_\_\_\_ was introduced by ...

- Solution: linear interpolation of two LMs
  - A global LM: using NLM
  - A local LM: using a fixed size cache, storing  $(\mathbf{h}_i, \mathbf{x}_{i+1})$

$$p(w_{t+1} \mid c_t) = (1 - \lambda)P_{\text{NLM}}(w_{t+1} \mid c_t) + \lambda P_{\text{cache}}(w_{t+1} \mid c_t)$$

$$P_{\text{cache}}(w_{t+1} \mid \underbrace{[\mathbf{h}_{1..t}], [\mathbf{x}_{1..t}]}_{c_t}) \propto \sum_{i=1}^{t-1} \mathbf{1}[w_{t+1} = x_{i+1}] \exp(\theta \langle \mathbf{h}_t, \mathbf{h}_i \rangle)$$

## Understand

$$P_{\text{cache}}(w_{t+1} \mid [\mathbf{h}_{1..t}], [\mathbf{x}_{1..t}]) \propto \sum_{i=1}^{t-1} \mathbf{1}[w_{t+1} = x_{i+1}] \underbrace{\exp(\theta \langle \mathbf{h}_t, \mathbf{h}_i \rangle)}_{\lambda_{t,i}}$$

- Programming-wise:
  - $\forall w \in V: w \rightarrow s_w$  by going through every item in the cache
  - then  $P(w|c_t) \propto s_w$
- Thinking of distributions: (let cache size be  $m$ )
  - Define distributions over  $V$ :  $\mathcal{D}_i := \mathbf{1}[w_{t+1} = x_{i+1}]$ 
    - We can represent it as a column vector.
  - $P(t \mid C) \propto \sum_i \lambda_{t,i} \cdot \mathcal{D}_i = \mathbf{D}\boldsymbol{\lambda}$ , where
    - $\mathcal{D}_i = \mathcal{D}_i$  (Note:  $\mathbf{D}$ 's shape is  $|V| \times m$ )
    - $\boldsymbol{\lambda} = \exp(\mathbf{H}\mathbf{h}_t)$ , where the  $i$ -th row of  $\mathbf{H}$  is  $\mathbf{h}_i$ .

# Neural LM with a Cache /3

## Advantages:

- Efficient way to adapt a LM to a new domain
- Can predict OOVs after seeing it once
- Captures long-range dependency

## Neural Language Model:

- Advantages (mainly compared with  $n$ -gram LMs):
  - Support much longer histories
  - Generalize better over contexts of similar words
  - Typically higher predictive accuracy
  - Integrates well with other DL-based methods (e.g., seq2seq model)
- Disadvantages:
  - Slower to train
  - Hard to interpret and understand the prediction results

- SPL 3ed, Chaps 7 and 9
- CS224n, Lecture 8: <http://web.stanford.edu/class/cs224n/lectures/lecture8.pdf>
- Edouard Grave, Armand Joulin, Nicolas Usunier: Improving Neural Language Models With A Continuous Cache. **ICLR 2017**. <https://openreview.net/forum?id=B184E5qee>